Ryan Yavari

San Francisco, CA (Open to Relocate) | 408-892-7723 | ryanyavari10@gmail.com | linkedin.com/in/ryan-yavari | ryanyavari.com

EDUCATION

Georgia Institute of Technology (Part-Time, Online) Atlanta, GA
Master of Science in Computer Science (Specialization in Artificial Intelligence). GPA: 3.67/4.0 *Expected Dec 2027*

San Jose State University San Jose, CA
Bachelor of Science in Computer Science. GPA: 3.67/4.0 *May 2025*

EXPERIENCE

Software Engineer Intern June 2025 - Jan. 2026
San Jose Mineta International Airport (SJC) *San Jose, CA*

- Engineered and deployed an LLM-powered customer service system with tool-augmented RAG pipeline serving 10,000+ travelers, reducing call volume by 17% and saving $11,000+ in annual support costs, live at flysanjose.com
- Implemented a production-grade RAG pipeline in Python with tool calling using LangChain, OpenAI embeddings, and Qdrant, enabling grounded, context-aware responses across 13,000+ chunked documents ingested from 667 pages of the SJC website
- Built a Redis-based semantic cache with hash-value change detection and automated refresh, reducing LLM inference costs by 25% and improving end-to-end response latency by 66%
- Architected a serverless deployment on Google Cloud Functions and Firebase with retry logic, graceful failure modes, and elastic scaling, maintaining 99.99% uptime across 6+ months in production
- Designed and ran a Python-based model selection evaluation suite across GPT variants, using an LLM-as-Judge on 200+ benchmark Q&A pairs, improving answer accuracy from 75% to 98% and reducing response latency from 5s to under 3s

PROJECTS

Optimal (AI Productivity Agent)| https://useoptimal.app June 2025 - Jan. 2026

- Engineered and launched an AI productivity platform used by 100+ users, built on a supervisor-orchestrated multi-agent system with stateful routing and structured agent handoffs
- Designed human-in-the-loop (HITL) approval gates where agents generate structured proposals for user review before persistence, enabling safe, auditable agentic actions with zero unreviewed writes
- Built offline evaluation benchmarks and regression suites to measure agent task completion, output quality, and planning correctness, cutting prompt regression incidents by 40% across 3 model update cycles
- Built full-stack infrastructure with Spring Boot, React Native (Expo), and secure token-based auth, designing a Supabase PostgreSQL schema with 17+ entities using EER modeling and normalization

PrepAgent (AI Interview Research Agent) | https://github.com/ryanyavari/prepagent Sept. 2026 - Oct. 2026

- Architected an MCP-native multi-agent system in Python using a LangGraph supervisor graph with 4 specialized agents, integrating Gmail, Google Calendar, and Google Drive via Google OAuth to autonomously detect interview invites and trigger a personalized research pipeline without manual user input
- Engineered a RAG pipeline with Pinecone vector store, Cohere Rerank, and OpenAI text-embedding-3-small to retrieve and rerank context across heterogeneous sources (resume, past briefings, company docs), with a RAGAS eval suite tracking answer relevance, faithfulness, and context precision across 20 ground truth briefings
- Automated a system to automate pre-interview research that previously took 45 minutes of manual review across resume, company docs, and past briefings, targeting sub-2-minute end-to-end generation

Quizzler (AI Quiz Generator) | https://github.com/ryanyavari/quizzler Mar. 2025, Updated Apr. 2026

- Instrumented a 5-stage RAG pipeline with LangSmith tracing across query expansion, embedding, reranking, and generation, capturing token usage, Cohere relevance scores, and per-span latency to surface retrieval quality and cost regressions
- Architected parallel upload processing to chunk and embed user documents into an ephemeral server-side vector store while concurrently generating structured MCQ quizzes, delivering grounded answers with glass-box citations from user-uploaded PDFs and text

SKILLS

Languages: Python, TypeScript, JavaScript, Java, SQL, Golang (Go), C, C++, HTML, CSS
AI & Machine Learning: RAG, Agentic Systems, Tool Calling, LangChain, LangGraph, MCP, OpenAI API, Anthropic API, Vector Databases, LLM Eval, RAGAS, Hallucination Detection, Observability, Function Calling, Prompt Engineering
Core Frameworks: FastAPI, Flask, Spring Boot, Node.js, Express, React, Next.js, React Native (Expo)
Databases & Cloud: PostgreSQL, Redis, GCP, AWS, Firebase, MongoDB, MySQL, Supabase, pgvector, Qdrant, Pinecone
Tools: Linux, Docker, Kubernetes, Git, CI/CD, RESTful APIs, System Design, Postman, LangSmith, LlamaIndex, Vercel AI